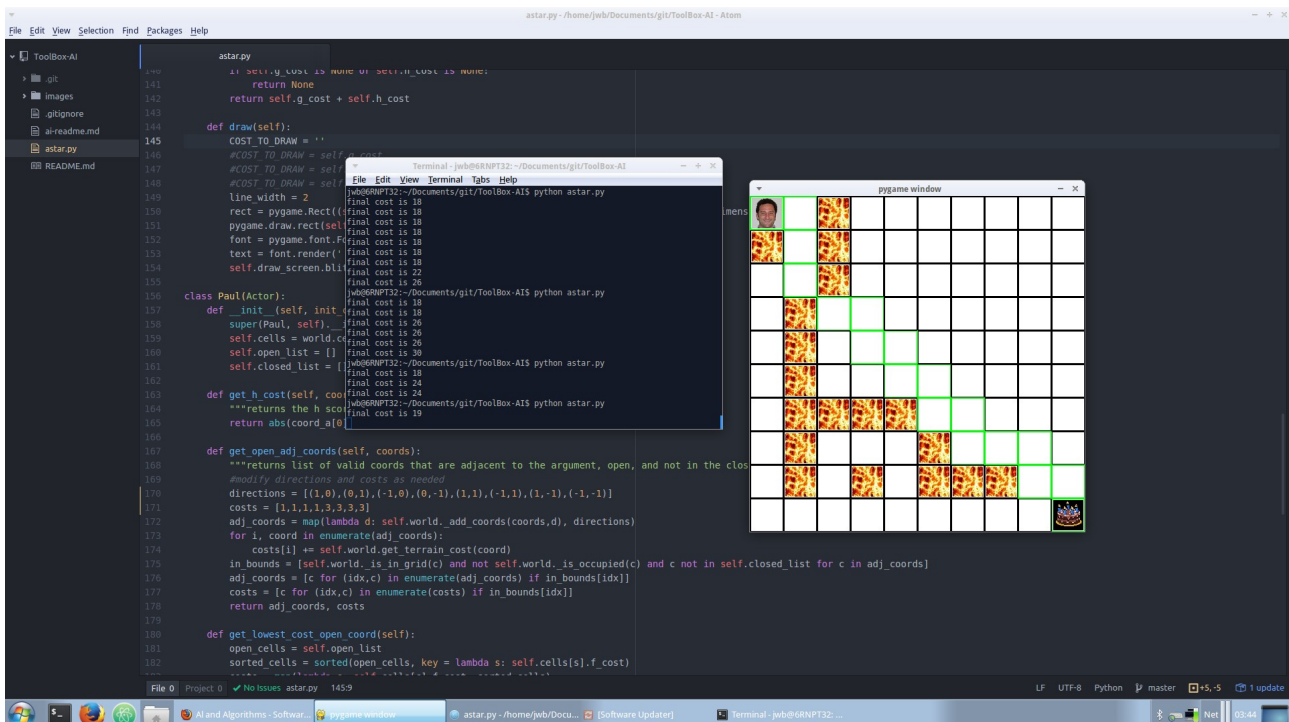
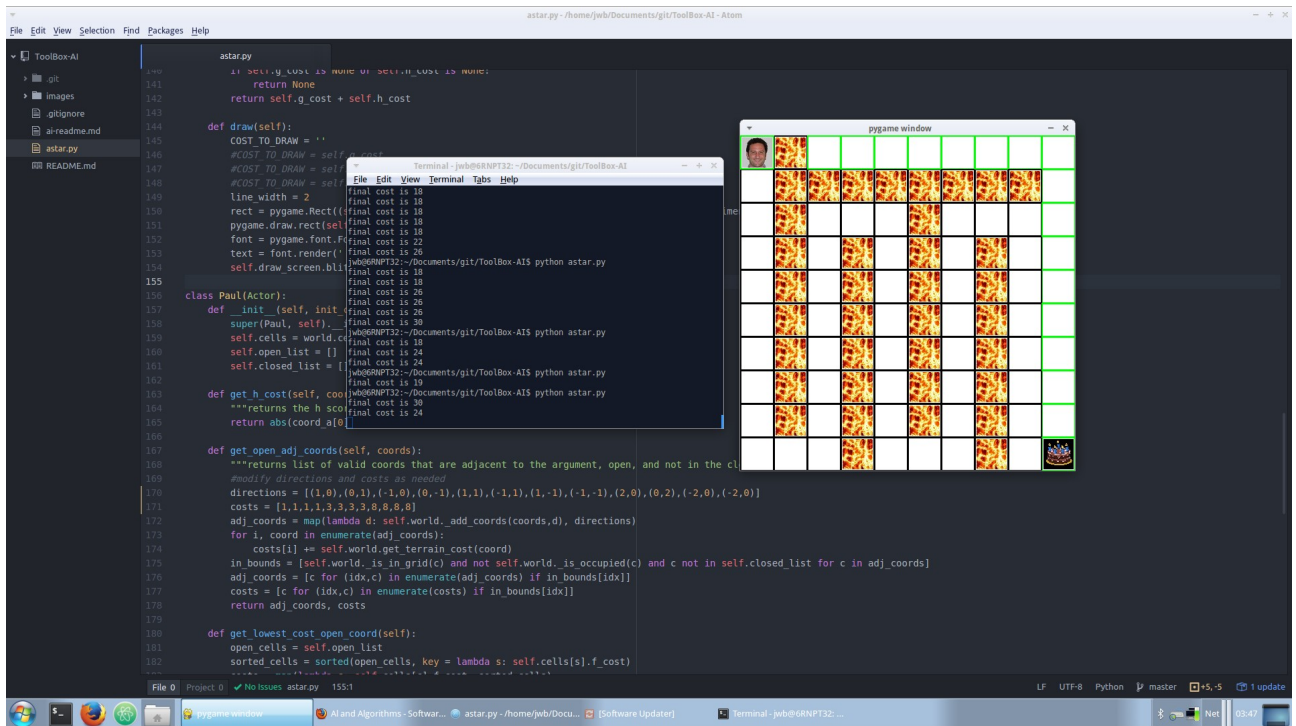


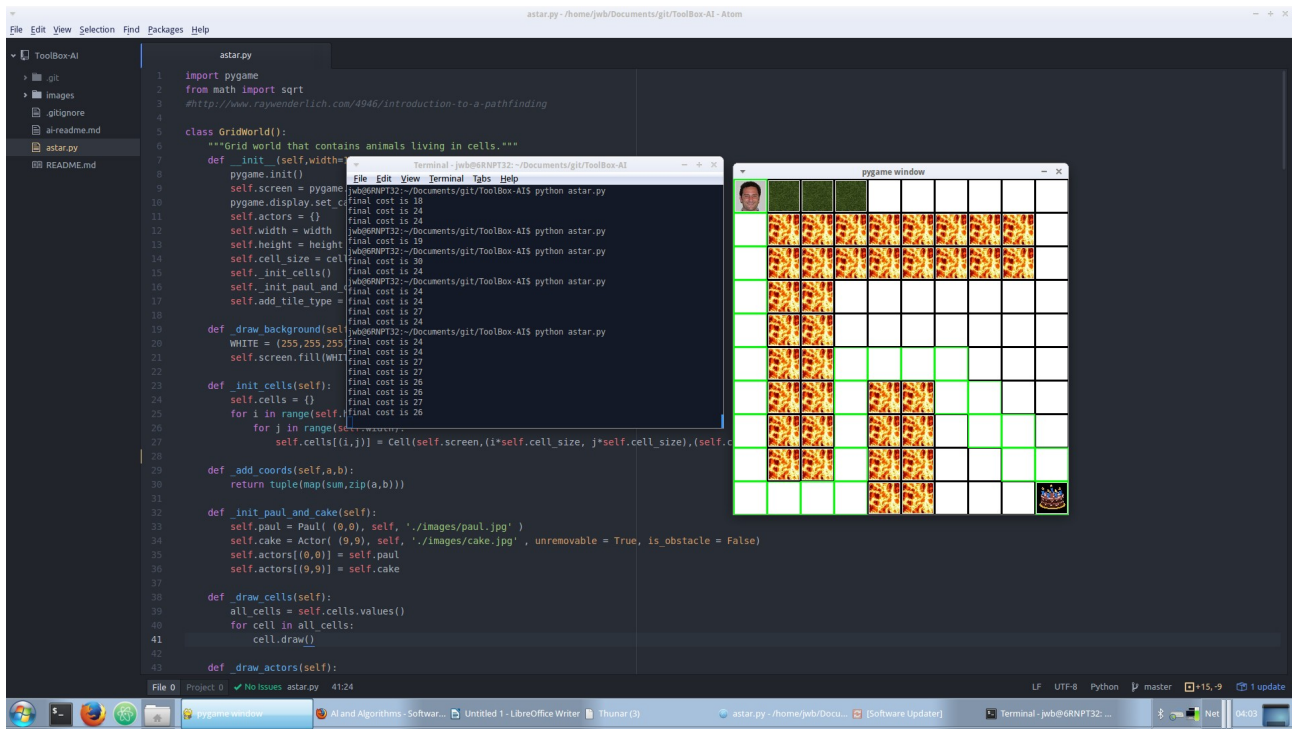
Here we see three different costs visualized: f, g and h. g is the cost that it takes to reach that node on our abstract graph, or grid location in this case. H is a prediction of how costly it will be from that node to the goal, or that grid square. F is just their sum.



Here we see diagonals implemented. Instead of taking a more costly edge route, paul can cut through a diagonal.



Here we see hops. Again, instead of a much more costly path in total, paul can just hop one square, even though this incurs a higher step cost.



Here we see swamp tiles. In the first picture, you can see paul is very averse to moving through swamp tiles, taking a longer but cheaper route. However, with enough costs incurred on all other options, the moving through swamps becomes the most efficient.